

[https://drive.google.com/file/d/1wYc6djaXSgazLsHSQn6PtKl\\_in5mGRck/view?usp=sharing](https://drive.google.com/file/d/1wYc6djaXSgazLsHSQn6PtKl_in5mGRck/view?usp=sharing)  
[\(https://drive.google.com/file/d/1wYc6djaXSgazLsHSQn6PtKl\\_in5mGRck/view?usp=sharing\)](https://drive.google.com/file/d/1wYc6djaXSgazLsHSQn6PtKl_in5mGRck/view?usp=sharing)

In [452]: `!gdown 1wYc6djaXSgazLsHSQn6PtKl_in5mGRck`

'gdown' is not recognized as an internal or external command,  
operable program or batch file.

```
In [128]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")

df=pd.read_csv("Jamboree.csv")
df.head()
```

Out[128]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	1	337	118	4	4.5	4.5	9.65	1	0.92
1	2	324	107	4	4.0	4.5	8.87	1	0.76
2	3	316	104	3	3.0	3.5	8.00	1	0.72
3	4	322	110	3	3.5	2.5	8.67	1	0.80
4	5	314	103	2	2.0	3.0	8.21	0	0.65

In [413]: `df.shape`

Out[413]: (500, 9)

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Serial No.            500 non-null   int64
1   GRE Score             500 non-null   int64
2   TOEFL Score           500 non-null   int64
3   University Rating     500 non-null   int64
4   SOP                   500 non-null   float64
5   LOR                   500 non-null   float64
6   CGPA                  500 non-null   float64
7   Research              500 non-null   int64
8   Chance of Admit       500 non-null   float64
dtypes: float64(4), int64(5)
memory usage: 35.3 KB
```

In [76]: `df["SOP"].value_counts()`

...

```
In [129]: df["University Rating"]=df["University Rating"].astype("category")
df["SOP"]=df["SOP"].astype("category")
df["LOR "]=df["LOR "].astype("category")
df["Research"]=df["Research"].astype("category")
```

```
In [236]: df.columns
```

```
Out[236]: Index(['Serial No.', 'GRE Score', 'TOEFL Score', 'University Rating', 'SOP',
               'LOR ', 'CGPA', 'Research', 'Chance of Admit '],
              dtype='object')
```

```
In [4]: df.isnull().sum().sum()
```

```
Out[4]: 0
```

There are no missing values found.

```
In [130]: df.drop("Serial No.",inplace=True,axis=1)
```

```
In [131]: df.describe(include="all")
```

```
Out[131]:
```

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
<b>count</b>	500.000000	500.000000	500.0	500.0	500.0	500.000000	500.0	500.00000
<b>unique</b>	NaN	NaN	5.0	9.0	9.0	NaN	2.0	NaN
<b>top</b>	NaN	NaN	3.0	4.0	3.0	NaN	1.0	NaN
<b>freq</b>	NaN	NaN	162.0	89.0	99.0	NaN	280.0	NaN
<b>mean</b>	316.472000	107.192000	NaN	NaN	NaN	8.576440	NaN	0.72174
<b>std</b>	11.295148	6.081868	NaN	NaN	NaN	0.604813	NaN	0.14114
<b>min</b>	290.000000	92.000000	NaN	NaN	NaN	6.800000	NaN	0.34000
<b>25%</b>	308.000000	103.000000	NaN	NaN	NaN	8.127500	NaN	0.63000
<b>50%</b>	317.000000	107.000000	NaN	NaN	NaN	8.560000	NaN	0.72000
<b>75%</b>	325.000000	112.000000	NaN	NaN	NaN	9.040000	NaN	0.82000
<b>max</b>	340.000000	120.000000	NaN	NaN	NaN	9.920000	NaN	0.97000

```
In [424]: df["LOR "].unique()
```

```
Out[424]: [4.5, 3.5, 2.5, 3.0, 4.0, 1.5, 2.0, 5.0, 1.0]
Categories (9, float64): [1.0, 1.5, 2.0, 2.5, ..., 3.5, 4.0, 4.5, 5.0]
```

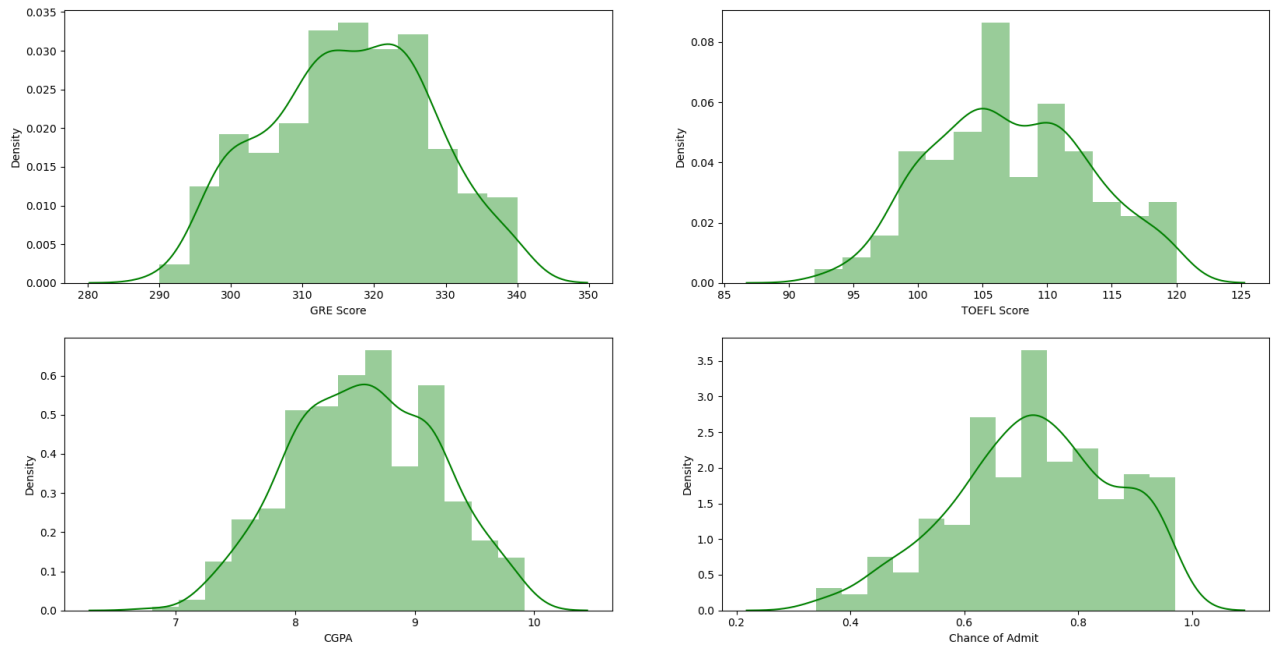
## UNIVARIATE ANALYSIS

```
In [396]: fig,axs=plt.subplots(nrows=2,ncols=2,figsize=(20,10))

cols=["GRE Score","TOEFL Score","CGPA","Chance of Admit "]
count=0

for i in range(2):
    for j in range(2):
        sns.distplot(df[cols[count]],ax=axs[i,j],color="g")
        count +=1

plt.show()
```



In [ ]: Above distribution plots shows that they are normally distributed

```
In [244]: df.info()
```





```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   GRE Score              500 non-null    int64
1   TOEFL Score            500 non-null    int64
2   University Rating      500 non-null    category
3   SOP                    500 non-null    category
4   LOR                    500 non-null    category
5   CGPA                   500 non-null    float64
6   Research               500 non-null    category
7   Chance of Admit        500 non-null    float64
dtypes: category(4), float64(2), int64(2)
memory usage: 18.8 KB
```

```
In [417]: fig,axs=plt.subplots(nrows=2,ncols=2,figsize=(20,20))

cols=["University Rating","SOP","LOR ","Research"]
count=0

for i in range(2):
    for j in range(2):
        sns.countplot(data=df,x=cols[count],ax=axs[i,j])
        count +=1


plt.show()
```



## Bivariate Analysis


```
In [6]: sns.barplot(data=df,x="Research",y="Chance of Admit ")
```

```
Out[6]: <Axes: xlabel='Research', ylabel='Chance of Admit ' >
```




```
In [64]: sns.barplot(data=df,x="University Rating",y="Chance of Admit ")
```

```
Out[64]: <Axes: xlabel='University Rating', ylabel='Chance of Admit ' >
```




```
In [71]: sns.barplot(data=df,x="SOP",y="Chance of Admit ")
```

```
Out[71]: <Axes: xlabel='SOP', ylabel='Chance of Admit ' >
```




```
In [73]: sns.barplot(data=df,x="LOR ",y="Chance of Admit ")
```

```
Out[73]: <Axes: xlabel='LOR ', ylabel='Chance of Admit ' >
```




```
In [65]: sns.barplot(data=df,x="University Rating",y="GRE Score")
```

```
Out[65]: <Axes: xlabel='University Rating', ylabel='GRE Score'>
```



```
In [74]: sns.barplot(data=df,x="University Rating",y="TOEFL Score")
```


```
Out[74]: <Axes: xlabel='University Rating', ylabel='TOEFL Score'>
```



```
In [ ]:
```


```
In [418]: sns.heatmap(df.corr(),annot=True)
```

```
Out[418]: <Axes: >
```



```
In [132]: sns.scatterplot(data=df,x="Chance of Admit ",y="GRE Score")
```


```
Out[132]: <Axes: xlabel='Chance of Admit ', ylabel='GRE Score'>
```






```
In [82]: sns.scatterplot(data=df,x="Chance of Admit ",y="TOEFL Score")
```

```
Out[82]: <Axes: xlabel='Chance of Admit ', ylabel='TOEFL Score'>
```

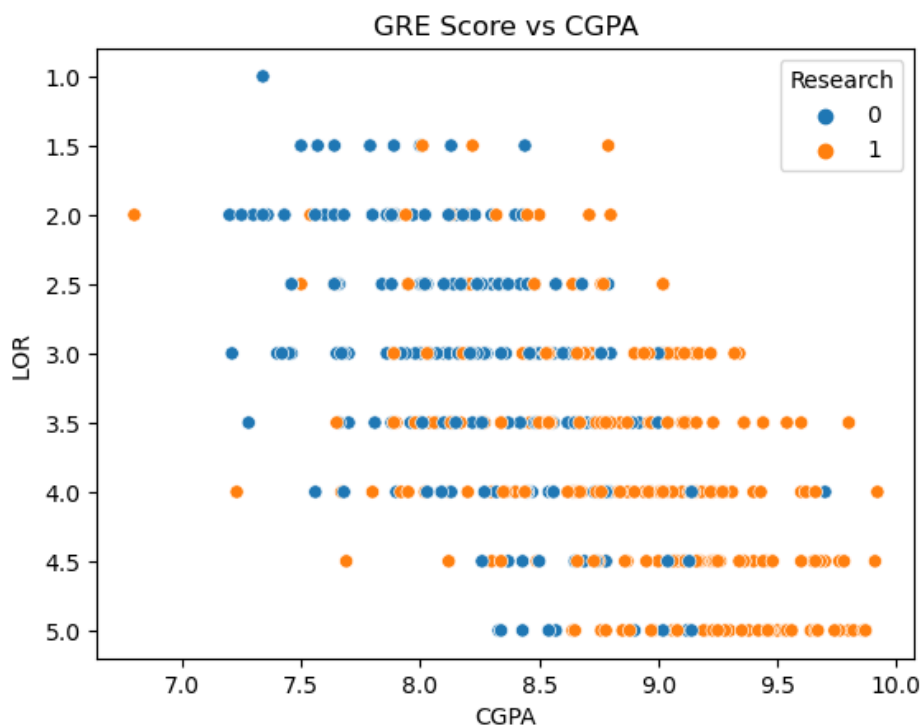


```
In [83]: sns.scatterplot(data=df,x="TOEFL Score",y="GRE Score")
```


```
Out[83]: <Axes: xlabel='TOEFL Score', ylabel='GRE Score'>
```




```
In [84]: fig = sns.scatterplot(x="CGPA", y="LOR ", data=df, hue="Research")
plt.title("GRE Score vs CGPA")
plt.show()
```




```
In [246]: fig = sns.scatterplot(x="TOEFL Score", y="LOR ", data=df, hue="Research")
plt.title("GRE Score vs CGPA")
plt.show()
```



```
In [86]: fig = sns.scatterplot(x="CGPA", y="SOP", data=df)
plt.title("GRE Score vs CGPA")
plt.show()
```



```
In [459]: fig = sns.scatterplot(x="TOEFL Score", y="SOP", data=df)
plt.title("GRE Score vs CGPA")
plt.show()
```

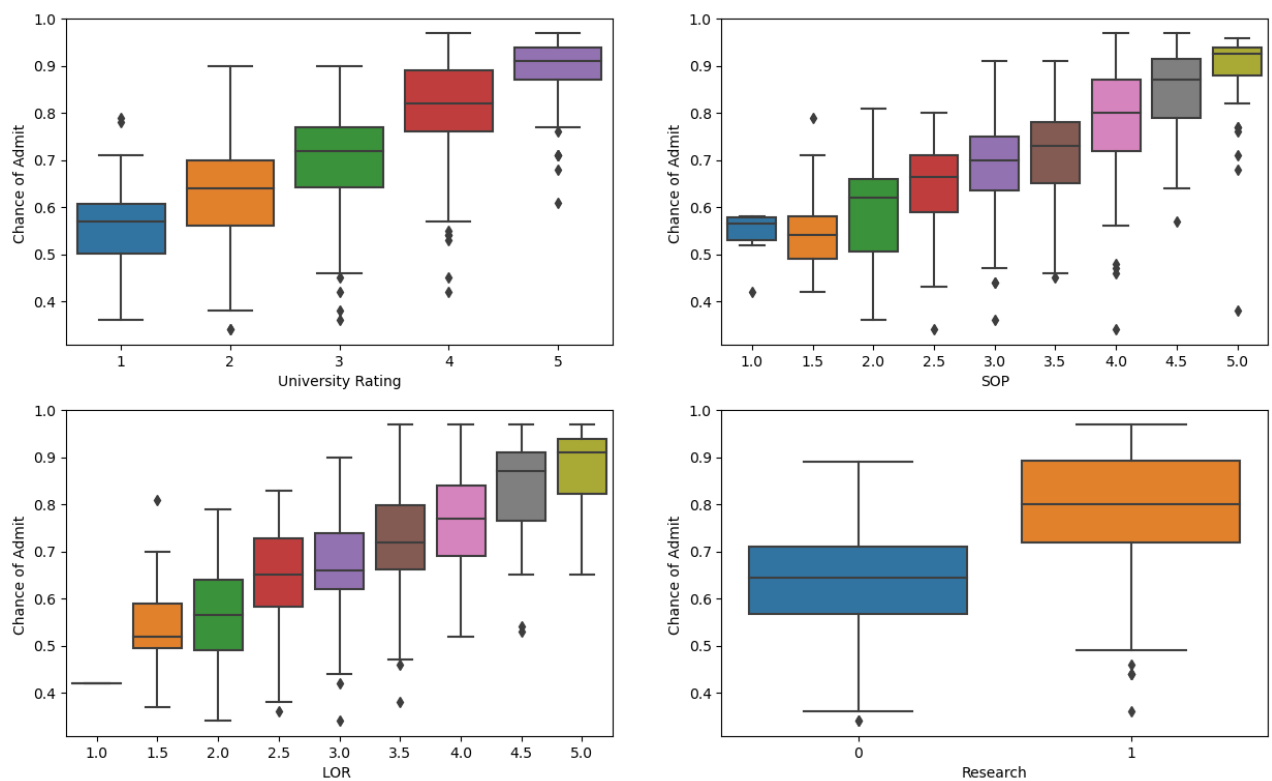


```
In [460]: df.dtypes
```

```
Out[460]: GRE Score          int64
TOEFL Score          int64
University Rating    category
SOP                  category
LOR                  category
CGPA                 float64
Research             category
Chance of Admit      float64
dtype: object
```

```
In [461]: fig,axs=plt.subplots(nrows=2,ncols=2,figsize=(15,9))
sns.boxplot(x="University Rating",y="Chance of Admit ",data=df,ax=axs[0,0])
sns.boxplot(x="SOP",y="Chance of Admit ",data=df,ax=axs[0,1])
sns.boxplot(x="LOR ",y="Chance of Admit ",data=df,ax=axs[1,0])
sns.boxplot(x="Research",y="Chance of Admit ",data=df,ax=axs[1,1])
```

```
Out[461]: <Axes: xlabel='Research', ylabel='Chance of Admit '>
```




```
In [462]: df.duplicated().sum()
```

```
Out[462]: 0
```

```
In [ ]: Treating outliers
```

```
In [463]: sns.distplot(df["GRE Score"],color="g")
```

```
Out[463]: <Axes: xlabel='GRE Score', ylabel='Density'>
```



```
In [133]: x=df.drop("Chance of Admit ",axis=1)
          y=df["Chance of Admit "]
```

```
In [134]: from sklearn.model_selection import train_test_split
          x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2)
```

```
In [135]: x_train
```

```
Out[135]:
```

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research
84	340	115	5	4.5	4.5	9.45	1
416	315	104	3	4.0	2.5	8.10	0
133	323	112	5	4.0	4.5	8.78	0
60	309	100	2	3.0	3.0	8.10	0
322	314	107	2	2.5	4.0	8.27	0
...	...	...	...	...	...	...	...
143	340	120	4	4.5	4.0	9.92	1
261	312	104	3	3.5	4.0	8.09	0
200	317	103	3	2.5	3.0	8.54	1
421	321	112	3	3.0	4.5	8.95	1
82	320	110	5	5.0	4.5	9.22	1

400 rows × 7 columns

```
In [137]: x_train.shape, y_train.shape, x_test.shape, y_test.shape
```

```
Out[137]: ((400, 7), (400,), (100, 7), (100,))
```

```
In [138]: from sklearn.preprocessing import StandardScaler
X_train_columns=x_train.columns
std=StandardScaler()
X_train_std=std.fit_transform(x_train)
```

```
In [139]: X_train_std
```

```
Out[139]: array([[ 2.08194   ,  1.31212211,  1.62731455, ...,  1.07999323,
                  1.49551735,  0.89543386],
                 [-0.1139676 , -0.49997371, -0.09927914, ..., -1.05596372,
                  -0.75727045, -1.11677706],
                 [ 0.58872283,  0.81791416,  1.62731455, ...,  1.07999323,
                  0.37746711, -1.11677706],
                 ...,
                 [ 0.061705  , -0.66470969, -0.09927914, ..., -0.52197448,
                  -0.0230285 ,  0.89543386],
                 [ 0.41305022,  0.81791416, -0.09927914, ...,  1.07999323,
                  0.6611515  ,  0.89543386],
                 [ 0.32521392,  0.48844219,  1.62731455, ...,  1.07999323,
                  1.11170906,  0.89543386]])
```

```
In [140]: from sklearn.metrics import mean_squared_error
from sklearn.linear_model import Lasso,Ridge,LinearRegression
from sklearn.metrics import accuracy_score
```

```
In [141]: X_train=pd.DataFrame(X_train_std, columns=X_train_columns)
```

```
In [142]: X_train
```

```
...
```

```
In [143]: lin = LinearRegression()
lin.fit(X_train,y_train.values)
y_pred3=lin.predict(std.transform(x_test))
lin.coef_
```

```
Out[143]: array([0.02417929, 0.01875231, 0.00598916, 0.00269915, 0.01675084,
                  0.06675081, 0.01017487])
```

```
In [144]: lin.intercept_
```

```
Out[144]: 0.7179750000000001
```

```
In [145]: mse = mean_squared_error(y_test,y_pred3)
print('Mean Squared Error = ',mse )

Mean Squared Error =  0.0033975693207486113
```

```
In [146]: reg = Lasso(alpha = 0.1)
reg.fit(X_train,y_train)
y_pred1=reg.predict(std.transform(x_test))
reg.coef_
```

```
Out[146]: array([0.          , 0.          , 0.          , 0.          , 0.          ,
                  0.02328122, 0.          ])
```

```
In [147]: reg.intercept_
```

```
Out[147]: 0.717975
```

```
In [148]: mse = mean_squared_error(y_test,y_pred1)
print('\nMean Squared Error = ',mse )
```

Mean Squared Error = 0.014711318628660916

```
In [149]: reg1 = Ridge(alpha = 1.0)
reg1.fit(X_train,y_train)
y_pred2=reg1.predict(std.transform(x_test))
reg1.coef_
```

```
Out[149]: array([0.02429928, 0.01888516, 0.00607765, 0.00284522, 0.01676942,
0.06619082, 0.0102024 ])
```

```
In [97]: reg1.intercept_
```

```
Out[97]: 0.722475
```

```
In [150]: mse = mean_squared_error(y_test,y_pred2)
print('\nMean Squared Error = ',mse )
```

Mean Squared Error = 0.003401805424686462

```
In [151]: import statsmodels.api as sm
X_train = sm.add_constant(X_train)
model = sm.OLS(y_train.values, X_train).fit()
print(model.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.819
Model:                  OLS      Adj. R-squared:          0.815
Method:                 Least Squares      F-statistic:        252.7
Date:                   Sun, 07 Jan 2024      Prob (F-statistic):    5.01e-141
Time:                   16:33:05      Log-Likelihood:        558.55
No. Observations:       400      AIC:                  -1101.
Df Residuals:           392      BIC:                  -1069.
Df Model:                7
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                0.7180        0.003    237.371     0.000        0.712        0.724
GRE Score            0.0242        0.006     3.889     0.000        0.012        0.036
TOEFL Score          0.0188        0.006     3.228     0.001        0.007        0.030
University Rating    0.0060        0.005     1.231     0.219       -0.004        0.016
SOP                  0.0027        0.005     0.532     0.595       -0.007        0.013
LOR                  0.0168        0.004     3.864     0.000        0.008        0.025
CGPA                 0.0668        0.006    10.310     0.000        0.054        0.079
Research             0.0102        0.004     2.758     0.006        0.003        0.017
=====
Omnibus:              87.250      Durbin-Watson:        2.129
Prob(Omnibus):         0.000      Jarque-Bera (JB):      199.415
Skew:                  -1.104      Prob(JB):              4.98e-44
Kurtosis:              5.663      Cond. No.              5.45
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

R-squared & Adj. R-squared both are almost same 82 % which is a good value because values greater than 0.5-1 are considered good fit model. Though it is noticed that SOP has low coefficient so it can be neglected.

```
In [152]: X_train_new = X_train.drop(columns="SOP")
```

```
In [153]: model1 = sm.OLS(y_train.values, X_train_new).fit()
print(model1.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.818
Model:                  OLS    Adj. R-squared:           0.816
Method:                 Least Squares    F-statistic:       295.3
Date:                   Sun, 07 Jan 2024    Prob (F-statistic): 3.22e-142
Time:                   16:33:13    Log-Likelihood:     558.41
No. Observations:       400    AIC:                 -1103.
Df Residuals:           393    BIC:                 -1075.
Df Model:                6
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.7180	0.003	237.588	0.000	0.712	0.724
GRE Score	0.0242	0.006	3.890	0.000	0.012	0.036
TOEFL Score	0.0189	0.006	3.260	0.001	0.008	0.030
University Rating	0.0070	0.004	1.551	0.122	-0.002	0.016
LOR	0.0175	0.004	4.263	0.000	0.009	0.026
CGPA	0.0674	0.006	10.611	0.000	0.055	0.080
Research	0.0101	0.004	2.754	0.006	0.003	0.017

```

=====
Omnibus:                 85.381    Durbin-Watson:           2.129
Prob(Omnibus):            0.000    Jarque-Bera (JB):         192.913
Skew:                     -1.086    Prob(JB):                 1.29e-42
Kurtosis:                 5.619    Cond. No.                  5.05
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

After removing it found that both R-squared & Adj. R-squared values are same with 82 % there is not much difference and still model remains fit.


## Testing the assumptions of LR model

### # 1)Linearity :

All the nuerical features like GRE TOEFL,CGPA are linear with Chance of Admit



```
In [176]: fig,(ax1,ax2,ax3) = plt.subplots(ncols=3,figsize=(12,2.5))
ax1.scatter(df["GRE Score"],df["Chance of Admit "])
ax1.set_title("GRE Score VS Chance of Admit ")
ax2.scatter(df["TOEFL Score"],df["Chance of Admit "])
ax2.set_title("TOEFL Score VS Chance of Admit ")
ax3.scatter(df["CGPA"],df["Chance of Admit "])
ax3.set_title("CGPA VS Chance of Admit ")
plt.show()
```



## # 2)Multicollinearity:

```
In [154]: from statsmodels.stats.outliers_influence import variance_inflation_factor
VIF=[]

for i in range(X_train_new.shape[1]):
    VIF.append(variance_inflation_factor(X_train_new,i))

pd.DataFrame({"VIF" : VIF},index=X_train_new.columns).T
```

Out[154]:

	const	GRE Score	TOEFL Score	University Rating	LOR	CGPA	Research
VIF	1.0	4.225817	3.681038	2.213631	1.842974	4.417958	1.486966

Since  $VIF < 5$ , there is low or no multicollinearity. Thus

## # Model performance evaluation"

Setting test dataset to match with dimension of train dataset.

```
In [156]: x_test_std=std.fit_transform(x_test)
```

```
In [158]: X_test=pd.DataFrame(x_test_std, columns=X_train_columns)
```

```
In [159]: X_test = sm.add_constant(X_test)
```

```
In [161]: X_test_del=list(set(X_test.columns).difference(set(X_train_new.columns)))
X_test_del
```

Out[161]: ['SOP']

```
In [162]: X_test_new=X_test.drop(columns=X_test_del)
```

```
In [191]: y_pred = model1.predict(X_test_new)

from sklearn.metrics import mean_squared_error,r2_score,mean_absolute_error

print('Mean Absolute Error ', mean_absolute_error(y_test.values,y_pred) )
print('Root Mean Square Error ', np.sqrt(mean_squared_error(y_test.values,y_pred) ))
print('R2 Score', r2_score(y_test.values,y_pred))

n=y_test.shape[0]
k = 1
Adj_r2score = 1-((1 - r2_score(y_test.values,y_pred)) * (n-1) / (n-k-1))
print('Adj_r2 Score', Adj_r2score)
```

```
Mean Absolute Error  0.04698336629208801
Root Mean Square Error  0.0613138027935518
R2 Score 0.8124412578774362
Adj_r2 Score 0.8105273931618998
```

MAE & RMSE values tends to zero and is very low which is good sign that model is performing well.  
R2 score and Adj\_r2score = 0.8 which tells that model capable of explaining 80% of variance of data.

```
In [197]: y_pred_t = model1.predict(X_train_new)

from sklearn.metrics import mean_squared_error,r2_score,mean_absolute_error

print('Mean Absolute Error ', mean_absolute_error(y_train.values,y_pred_t) )
print('Root Mean Square Error ', np.sqrt(mean_squared_error(y_train.values,y_pred_t) ))
print('R2 Score', r2_score(y_train.values,y_pred_t))

n=y_train.shape[0]
k = 1
Adj_r2score = 1-((1 - r2_score(y_train.values,y_pred_t)) * (n-1) / (n-k-1))
print('Adj_r2 Score', Adj_r2score)
```

```
Mean Absolute Error  0.04319820422265458
Root Mean Square Error  0.05990745344713923
R2 Score 0.8184594131773991
Adj_r2 Score 0.8180032810497042
```

MAE & RMSE values tends to zero and is very low which is good sign that model is performing well.  
R2 score and Adj\_r2score = 0.8 which tells that model capable of explaining 80% of variance of data.

### # 3)Normality of Residuals:


```
In [ ]: residuals = y_test - y_pred
```

```
In [165]: np.mean(residuals)
```

```
Out[165]: 0.010404491966023461
```

```
In [166]: sns.distplot(residuals,kde=True)
```


```
Out[166]: <Axes: xlabel='Chance of Admit ', ylabel='Density'>
```



```
In [295]: x_test_new.shape
```

```
Out[295]: (100, 6)
```


```
In [167]: import scipy as sp
fig,ax=plt.subplots(figsize=(6,4))
sp.stats.probplot(residuals,plot=ax,fit=True)
plt.show()
```



**# Test for Homoscedasticity**

```
In [181]: p = sns.scatterplot(x=pred,y=residuals)
plt.xlabel('predicted values')
plt.ylabel('Residuals')
plt.ylim(-0.4,0.4)
plt.xlim(0,1)
```

Out[181]: (0.0, 1.0)




when you imagine a straightline passing at 0 of residual y-axis , you can see the plot is uniform or same on both sides of the plot.

## # Autocorrelation check

```
In [182]: plt.plot(residuals)
```

Out[182]: [



No autocorrelation , that is no pattern found in residual data.

In [ ]: